


The big picture

Every building block, and how they connect. Knowledge is prepared once (top row); answers are produced on demand (middle row); patient data stays in its own per-practitioner vault. Follow the arrows.



How the clinic teaches its AI

Before the assistant can help anyone, the library fills its own shelves. An admin uploads a source and everything else happens automatically, in one write.

1

Upload

An admin drops in a source — a PDF, pasted text, or a scraped web page — via POST /api/sources (or the chunked upload flow for anything large).

2

Extract & split

Text is pulled per page from the file, then split into passages — the citable unit — with page numbers preserved for later citation.

3

Read & grade

The Reader titles it, summarises it, tags its topics (reusing an existing shelf wherever one fits), and proposes a 1–10 reliability grade from the strength of the evidence. An admin can revise the grade or topics at any time.

4

Embed & extract the graph

The Indexer embeds every passage; the Cartographer pulls out the clinical entities and relationships that passage actually states — one failed passage never blanks the rest of the document.

engineers call this: embeddings, held in Neo4j's own vector index — semantic search, not just keyword match.

5

One atomic write

Document, Chunks, and Entities land together in Neo4j — or not at all. The source is immediately searchable (seed search) and traversable (graph expansion) the moment it lands; it never sits as a disconnected island.

🔍 The admin Library shows every source's grade, topics and ingest date, plus the graph's running totals (documents, chunks, entities, typed relationships). An exact re-upload of the same content is blocked automatically (content-hash check) rather than silently duplicated. Each practitioner sets their own reliability floor — nothing below their grade threshold reaches an answer for them.

How a question gets answered

A real moment from the consultation room — this is the heart of the product.

1

The question

During a session, the practitioner opens the client's file and asks the AI team a question, scoped to that patient.

2

The Librarian listens

It reads the question plus the patient's context and forms the seed search query. It doesn't answer — it gathers.

3

Two sources today

It pulls from the clinic's own graded knowledge base ① (seed search, then LLM-judged graph traversal hop by hop) and the patient's own record ② — labs, history, past sessions, from their practitioner's Postgres vault. A third source — live PubMed lookup — is planned, not built yet.

4

The Specialist writes the answer

A senior model studies everything accumulated — fresh, every time — and writes the answer using only that material, with inline citations showing exactly where each point comes from.

engineers call this: retrieval-augmented generation (RAG) — grounded through a knowledge graph, not a flat vector search.

5

The anti-hallucination check

Before the answer is shown, HHEM scores every sentence against its cited source. A weak verdict sends the draft back to the Specialist once, with the unsupported lines named — a clean pass, a bounded retry, or an honestly flagged answer.

checker: HHEM-2.1-Open (vectara/hallucination_evaluation_model), open-source, run locally — not another LLM grading its own kind's work.

OPTIONAL

6

The answer

Back in the chat within seconds, with clickable citations back to the exact passage. The full turn — question, answer, traversal path, check verdict — is saved to the session automatically.

7

It remembers, on request

The practitioner can generate a short session summary at any point — saved straight to the client's record, and explicitly flagged if any answer in that session wasn't a clean check pass.

Meet the AI team

Not one do-everything model, but five specialised roles — each on the model actually suited to that job — plus a sixth, non-LLM teammate that checks every answer before a practitioner sees it. Every role can be independently repointed to a different provider without touching business logic.



INGESTION

The Reader

Reads a new source, writes its title, summary, topics and a suggested 1–10 reliability grade.

READER_MODEL · Qwen/Qwen3-30B-A3B-Instruct-2507



INGESTION

The Cartographer (Graph-builder)

Pulls the clinical entities and relationships out of each passage and links them, so related material across the library can find each other.

GRAPH_BUILDER_MODEL · Qwen/Qwen3-235B-A22B-Instruct-2507



INGESTION

The Indexer (Embedder)

Embeds every knowledge-base chunk once, at ingest time, for seed (vector + full-text) search at question time.

EMBEDDER_MODEL · Qwen/Qwen3-Embedding-8B



RETRIEVAL

The Librarian (Answer Engine)

Forms the seed query, then judges which candidates are worth expanding at every hop of graph traversal — never told the grade, so relevance and trust stay decided separately.

RETRIEVAL_MODEL · nvidia/Nemotron-3_5-Lightning



ANSWER

The Specialist (Reasoner)

The senior model — studies the accumulated, graded evidence and the patient's own data, and writes the grounded answer, every claim cited.

REASONER_MODEL · moonshotai/Kimi-K3



OPTIONAL

QUALITY

The Checker

Independently scores each answer sentence against its cited evidence before a practitioner sees it — not a chat model, a dedicated classifier, run locally.

CHECKER_MODEL · vectara/hallucination_evaluation_model (HHEM-2.1-Open)

💡 All five chat-model roles currently run on Nebius's OpenAI-compatible token factory (one shared key, one bill) — but nothing in ingestion, retrieval or reasoning code constructs a client or knows a base_url. Repointing one role to a different provider is a config change, not a code change.