



1. Ingestion & Upload Pipeline

One source (PDF, pasted text, or scraped page) → Document + Chunks + Entities in Neo4j


INTAKE




READING & GRADING



PER-PASSAGE PROCESSING (PARALLEL OVER EVERY CHUNK)




WRITE



■ LLM / AI-team call ■ Graph database write ■ User / admin action

2. Query → Answer → Hallucination Check

A clinician's question → a grounded, cited answer, checked for unsupported claims before it's shown



3. Knowledge Graph Schema (Neo4j)

What ingestion writes, and what retrieval traverses

