



1. Ingestion & Upload Pipeline

One source (PDF, pasted text, or scraped page) → Document + Chunks + Entities in Neo4j


INTAKE




READING & GRADING



PER-PASSAGE PROCESSING (PARALLEL OVER EVERY CHUNK)




WRITE



■ LLM / AI-team call ■ Graph database write ■ User / admin action


2. Query → Answer → Hallucination Check

A clinician's question → a grounded, cited answer, checked for unsupported claims before it's shown



3. Knowledge Graph Schema (Neo4j)

What ingestion writes, and what retrieval traverses



Patient context

Lives outside Neo4j (Postgres core + per-practitioner vault schema, not the graph) — conditions/meds/history feed the seed query and traversal judging, but are never written into the KG.

Practitioner scoping

Every Document carries a grade; a practitioner's own source_weights (Postgres vault schema) can raise or lower a source's effective weight without touching the graph itself.

■ Written by ingestion (page 1) ■ Written by the Graph Builder, traversed by retrieval (page 2) ■ Facet / browse structure